

Introduction

Final Project Submission

- Student Name: Adam Marianacci
- Student Pace: Flex
- Scheduled project review date/time: TBD
- Instructor Name: Morgan Jones

Business Understanding

Coldwell Banker wants to buy undervalued homes in certain zipcodes to minimally renovate and sell for a profit. It is my job to find homes that are being listed significantly lower than the median price and look into what variables are causing these lowered prices so that I can properly advise Coldwell Banker on which homes to target.

Data Understanding

This project uses the King County House Sales dataset. It contains over 30,000 entries of data related to the sale price of houses, number of bedrooms, bathrooms and floors, square footage, addresses, and more. The main dataframe used in this project only contained roughly 14,000 entries. A limitation of the data was that it is fairly small since we are dealing with predictive modeling. Also the features of the data were strongly correlated with each other (multicollinearity) rather than the target variable. The dataset is suitable for this project because it has information to reveal which homes are truly being undervalued in certain zipcodes. After analyzing the price, location, and other various specifications and amenities of homes I will be able to make informed recommendations to the real estate agency.

Data Preperation

I started by importing the necessary libraries.

```
In [1]: 1 import pandas as pd
        2 import re
        3 import numpy as np
        4 from sklearn import linear_model
        5 from sklearn.feature_selection import RFE
        6 from sklearn.tree import DecisionTreeClassifier
        7 from sklearn.model_selection import train_test_split
        8 from sklearn.linear_model import LinearRegression
        9 from sklearn.preprocessing import OneHotEncoder
       10 from sklearn.metrics import r2_score, mean_squared_error, mean_abso
       11 import seaborn as sns
       12 import matplotlib.pyplot as plt
       13 import statsmodels.api as sm
       14 import scipy.stats as stats
       15 import warnings
       16 warnings.filterwarnings('ignore')
```

I did not want any information in the dataframe to be truncated. I searched pandas output truncated in google and found this [solution \(https://stackoverflow.com/questions/25351968/how-can-i-display-full-non-truncated-dataframe-information-in-html-when-conver\)](https://stackoverflow.com/questions/25351968/how-can-i-display-full-non-truncated-dataframe-information-in-html-when-conver).

```
In [2]: 1 pd.set_option("max_columns", None) # show all cols
        2 pd.set_option('max_colwidth', None) # show full width of showing co
        3 pd.set_option("expand_frame_repr", False) # print cols side by side
        4 pd.set_option('display.max_rows', 500) #show all rows
        5 pd.options.display.float_format = '{:.2f}'.format #surpressing scie
```

```
In [3]: 1 # Loading the dataframe
        2 kc_house_data_df = pd.read_csv('data/kc_house_data.csv', index_col=
```

```
In [4]: 1 # Previewing the dataframe
        2 kc_house_data_df.head()
```

Out[4]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
id								
7399300360	5/24/2022	675000.00	4	1.00	1180	7140	1.00	NC
8910500230	12/13/2021	920000.00	5	2.50	2770	6703	1.00	NC
1180000275	9/29/2021	311000.00	6	2.00	2880	6156	1.00	NC
1604601802	12/14/2021	775000.00	3	3.00	2160	1400	2.00	NC
8562780790	8/24/2021	592500.00	2	2.00	1120	758	2.00	NC

There are 24 columns and 30155 entries. There is only a handful of missing values in the "heat_source" and "sewer_system" columns. The data types consist of objects, integers, and floats.

```
In [5]: 1 # Gathering information about the datatypes within the dataframe, as  
2        kc_house_data_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 30155 entries, 7399300360 to 9557800100  
Data columns (total 24 columns):  
#   Column                Non-Null Count  Dtype  
---  -  
0   date                   30155 non-null  object  
1   price                   30155 non-null  float64  
2   bedrooms               30155 non-null  int64  
3   bathrooms              30155 non-null  float64  
4   sqft_living             30155 non-null  int64  
5   sqft_lot                30155 non-null  int64  
6   floors                 30155 non-null  float64  
7   waterfront             30155 non-null  object  
8   greenbelt              30155 non-null  object  
9   nuisance               30155 non-null  object  
10  view                   30155 non-null  object  
11  condition              30155 non-null  object  
12  grade                  30155 non-null  object  
13  heat_source            30123 non-null  object  
14  sewer_system           30141 non-null  object  
15  sqft_above             30155 non-null  int64  
16  sqft_basement          30155 non-null  int64  
17  sqft_garage            30155 non-null  int64  
18  sqft_patio             30155 non-null  int64  
19  yr_built               30155 non-null  int64  
20  yr_renovated           30155 non-null  int64  
21  address                30155 non-null  object  
22  lat                    30155 non-null  float64  
23  long                   30155 non-null  float64  
dtypes: float64(5), int64(9), object(10)  
memory usage: 5.8+ MB
```

I wanted to confirm and see clearly how many missing values there were in each column.

```
In [6]: 1 kc_house_data_df.isnull().sum()
```

```
Out[6]: date                0
price                0
bedrooms            0
bathrooms           0
sqft_living         0
sqft_lot            0
floors              0
waterfront          0
greenbelt           0
nuisance             0
view                0
condition           0
grade               0
heat_source         32
sewer_system        14
sqft_above           0
sqft_basement        0
sqft_garage          0
sqft_patio           0
...             ^
```

Since there were only 32 missing values from "heat_source" and 14 missing values from "sewer_system" out of 30,155 entries, I decided to remove the rows that contained missing values. I typed remove rows that contain missing values in a pandas df into google and found this [solution \(https://stackoverflow.com/questions/13413590/how-to-drop-rows-of-pandas-dataframe-whose-value-in-a-certain-column-is-nan\)](https://stackoverflow.com/questions/13413590/how-to-drop-rows-of-pandas-dataframe-whose-value-in-a-certain-column-is-nan).

```
In [7]: 1 kc_house_data_df = kc_house_data_df.dropna(subset = ["heat_source",
```

Checking to see if my code worked, there are no more missing values.

In [8]: 1 kc_house_data_df.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 30111 entries, 7399300360 to 9557800100
Data columns (total 24 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date                  30111 non-null  object
1   price                 30111 non-null  float64
2   bedrooms              30111 non-null  int64
3   bathrooms             30111 non-null  float64
4   sqft_living           30111 non-null  int64
5   sqft_lot              30111 non-null  int64
6   floors                30111 non-null  float64
7   waterfront            30111 non-null  object
8   greenbelt            30111 non-null  object
9   nuisance              30111 non-null  object
10  view                  30111 non-null  object
11  condition             30111 non-null  object
12  grade                 30111 non-null  object
13  heat_source           30111 non-null  object
14  sewer_system          30111 non-null  object
15  sqft_above            30111 non-null  int64
16  sqft_basement         30111 non-null  int64
17  sqft_garage           30111 non-null  int64
18  sqft_patio            30111 non-null  int64
19  yr_built              30111 non-null  int64
20  yr_renovated          30111 non-null  int64
21  address               30111 non-null  object
22  lat                   30111 non-null  float64
23  long                  30111 non-null  float64
dtypes: float64(5), int64(9), object(10)
memory usage: 5.7+ MB
```

"Lat", "Long", did not seem relevant since I am looking into zipcodes. I dropped them from the dataframe to simplify and condense the df.

In [9]: 1 kc_house_data_df.drop(['lat', 'long'], axis=1, inplace=True)

I wanted to extract zipcodes from addresses so I searched extracting zipcodes from addresses in python and found this [solution \(https://stackoverflow.com/questions/66384707/extracting-zip-code-from-a-string-with-full-address\)](https://stackoverflow.com/questions/66384707/extracting-zip-code-from-a-string-with-full-address).

```
In [10]: 1 def extract_zipcode(address):
2         zipcode = re.search(r'\b\d{5}\b', address)
3         if zipcode:
4             return zipcode.group(0)
5         else:
6             return None
7
8         # Create the "Zip Code" column by applying the function to the "address" column
9         kc_house_data_df['Zip Code'] = kc_house_data_df['address'].apply(extract_zipcode)
```

```
In [11]: 1 kc_house_data_df.head()
```

```
Out[11]:
```

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
id								
7399300360	5/24/2022	675000.00	4	1.00	1180	7140	1.00	NC
8910500230	12/13/2021	920000.00	5	2.50	2770	6703	1.00	NC
1180000275	9/29/2021	311000.00	6	2.00	2880	6156	1.00	NC
1604601802	12/14/2021	775000.00	3	3.00	2160	1400	2.00	NC
8562780790	8/24/2021	592500.00	2	2.00	1120	758	2.00	NC

After previewing the dataset again I noticed that it was interpreting 5 digit street addresses as zipcodes. To fix this problem I decided to choose the 5 digit numbers after the word "Washington".

```
In [12]: 1 # Function to extract ZIP code after "Washington"
2 def extract_zipcode(address):
3     match = re.search(r'Washington (\d{5})', address)
4     if match:
5         return match.group(1)
6     return None
7
8 kc_house_data_df['Zip Code'] = kc_house_data_df['address'].apply(ex
```

```
In [13]: 1 kc_house_data_df.head()
```

Out[13]:

	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
id								
7399300360	5/24/2022	675000.00	4	1.00	1180	7140	1.00	NC
8910500230	12/13/2021	920000.00	5	2.50	2770	6703	1.00	NC
1180000275	9/29/2021	311000.00	6	2.00	2880	6156	1.00	NC
1604601802	12/14/2021	775000.00	3	3.00	2160	1400	2.00	NC
8562780790	8/24/2021	592500.00	2	2.00	1120	758	2.00	NC

Using the `nunique` function to find the number of zipcodes in the dataset which is 92.

```
In [14]: 1 unique_zipcodes = kc_house_data_df['Zip Code'].nunique()
```

```
In [15]: 1 unique_zipcodes
```

```
Out[15]: 92
```

Data Analysis

I started by running a `df.describe()` on the dataset to learn more about the descriptive statistics. I see that the mean of price is about 1.1 million and the median is 860,000, which means we do not have a normal distribution and it is right skewed. The majority of the data is located on the left side of the graph.

```
In [16]: 1 kc_house_data_df.describe()
```

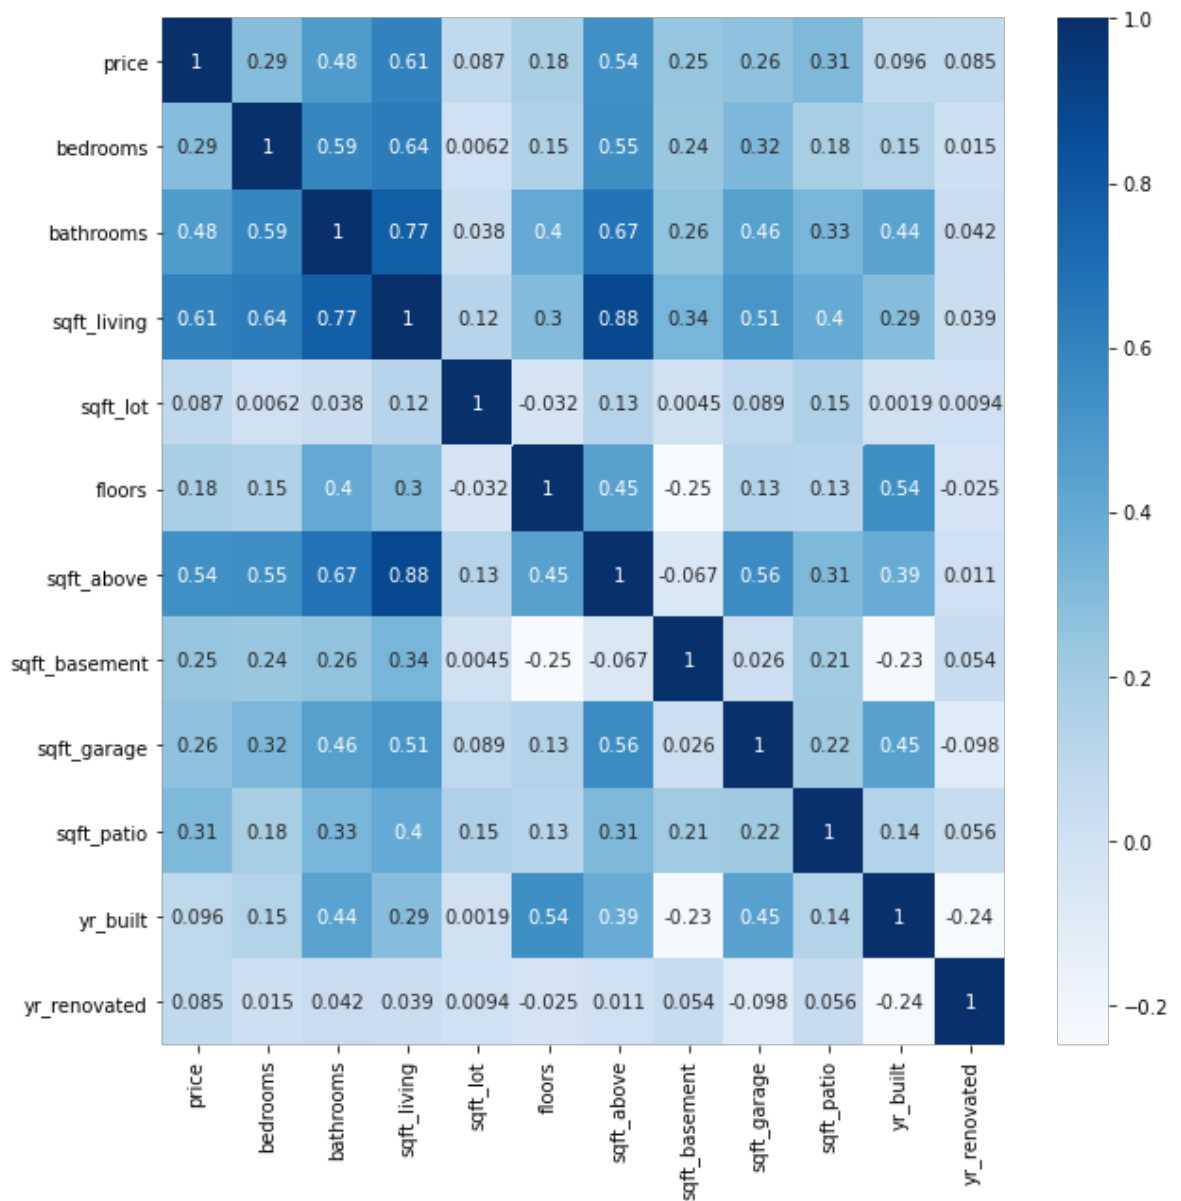
```
Out[16]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sqft_above	sqft_b
count	30111.00	30111.00	30111.00	30111.00	30111.00	30111.00	30111.00	30111.00
mean	1108970.65	3.42	2.34	2113.34	16648.80	1.54	1810.39	1810.39
std	896515.83	0.98	0.89	973.45	59933.03	0.57	877.73	877.73
min	27360.00	0.00	0.00	3.00	402.00	1.00	2.00	2.00
25%	649236.00	3.00	2.00	1420.00	4850.00	1.00	1180.00	1180.00
50%	860000.00	3.00	2.50	1920.00	7477.00	1.50	1560.00	1560.00
75%	1300000.00	4.00	3.00	2620.00	10568.00	2.00	2270.00	2270.00
max	30750000.00	13.00	10.50	15360.00	3253932.00	4.00	12660.00	12660.00

I used a heatmap to see how the features were correlated with each other and found that price was correlated the highest with `sqft_living`.

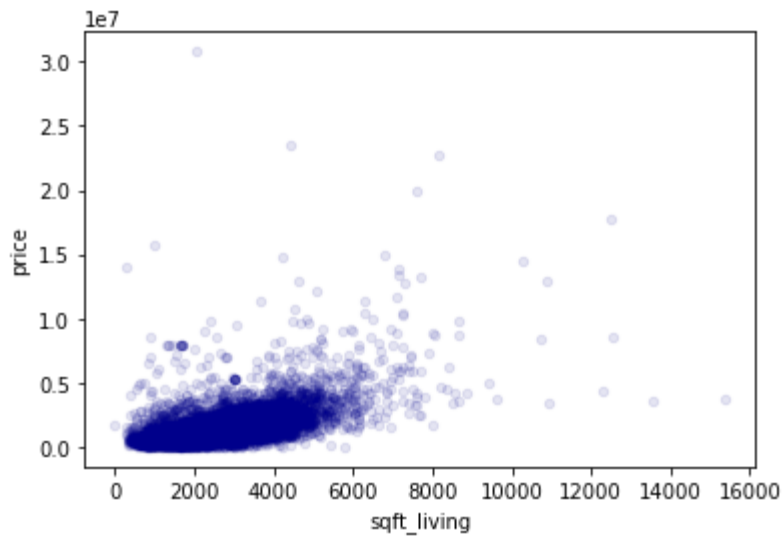
```
In [17]: 1 fig, ax = plt.subplots(figsize=(10,10))
          2 cor = kc_house_data_df.corr()
          3 sns.heatmap(cor, cmap="Blues", annot=True)
```

Out[17]: <AxesSubplot:>



```
In [18]: 1 kc_house_data_df.plot.scatter(  
2         x='sqft_living',  
3         y='price',  
4         c='DarkBlue',  
5         alpha = 0.1)
```

```
Out[18]: <AxesSubplot:xlabel='sqft_living', ylabel='price'>
```



I wanted to see how many zipcodes there were in the dataset and how many each contained so I ran a `value_counts` on "Zip Codes"

```
In [19]: 1 kc_house_data_df['Zip Code'].value_counts()
```

```
Out[19]: 98042    992
          98038    857
          98103    761
          98115    760
          98117    748
          98023    695
          98034    689
          98058    682
          98133    664
          98001    623
          98092    609
          98033    608
          98118    600
          98059    583
          98052    568
          98106    543
          98031    530
          98006    526
          98056    520
          98155    515
          98125    489
          98045    466
          98107    458
          98022    457
          98126    440
          98003    440
          98122    440
          98144    439
          98146    428
          98074    427
          98198    425
          98075    411
          98199    410
          98008    397
          98053    393
          98116    387
          98178    383
          98168    383
          98027    382
          98030    377
          98002    363
          98028    349
          98040    349
          98166    344
          98072    343
          98105    330
          98004    322
          98029    321
          98010    314
          98112    307
          98108    305
          98177    301
          98136    299
          98065    292
```

98011	261
98077	246
98019	245
98119	241
98055	218
98188	200
98070	191
98032	184
98005	178
98014	156
98007	154
98102	144
98109	129
98057	127
98024	114
98148	94
98047	77
98051	66
98039	59
98354	23
98288	16
98272	6
98271	5
98223	4
98224	3
98251	3
98338	2
98663	2
98050	2
98372	2
99202	1
98422	1
98296	1
98387	1
99403	1
99203	1
98270	1
99223	1

Name: Zip Code, dtype: int64

After reviewing the counts, I decided to drop the zipcodes that had under 100 houses listed to sale. I decided to keep the '98149' zipcode with a count of 94 because it was close.

```
In [20]: 1 # Count the occurrences of each ZIP code
          2 zipcode_counts = kc_house_data_df['Zip Code'].value_counts()
          3 kc_house_data_df = kc_house_data_df[kc_house_data_df['Zip Code'].is
          4
```

Checking to see if the Zip Codes below a count of 94 were dropped, they were!

```
In [21]: 1 kc_house_data_df['Zip Code'].value_counts()
```

```
Out[21]: 98042    992
          98038    857
          98103    761
          98115    760
          98117    748
          98023    695
          98034    689
          98058    682
          98133    664
          98001    623
          98092    609
          98033    608
          98118    600
          98059    583
          98052    568
          98106    543
          98031    530
          98006    526
          98056    520
          98155    515
          98125    489
          98045    466
          98107    458
          98022    457
          98126    440
          98122    440
          98003    440
          98144    439
          98146    428
          98074    427
          98198    425
          98075    411
          98199    410
          98008    397
          98053    393
          98116    387
          98168    383
          98178    383
          98027    382
          98030    377
          98002    363
          98040    349
          98028    349
          98166    344
          98072    343
          98105    330
          98004    322
          98029    321
          98010    314
          98112    307
          98108    305
          98177    301
          98136    299
          98065    292
```

```

98011    261
98077    246
98019    245
98119    241
98055    218
98188    200
98070    191
98032    184
98005    178
98014    156
98007    154
98102    144
98109    129
98057    127
98024    114
98148     94
Name: Zip Code, dtype: int64

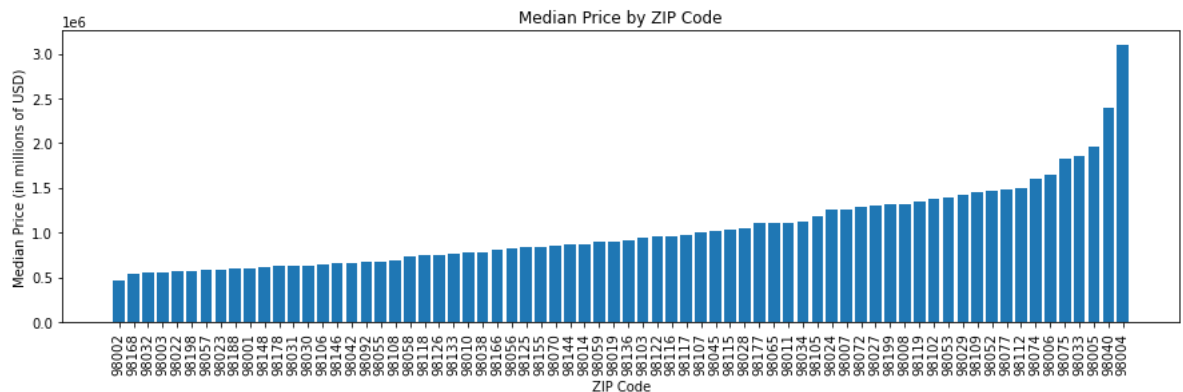
```

I wanted to create a graph of median prices of homes by zipcode to get an idea of which zipcodes had the highest median price. To do this I used the `sort_values` function along with the `median` function.

```

In [22]: 1 # Group data by ZIP code and calculate the median price from lowest
          2 median_prices = kc_house_data_df.groupby('Zip Code')['price'].media
          3
          4 # Create a bar chart to visualize the median prices by ZIP code
          5 plt.figure(figsize=(15, 4))
          6 plt.bar(median_prices['Zip Code'], median_prices['price'])
          7 plt.xlabel('ZIP Code')
          8 plt.ylabel('Median Price (in millions of USD)')
          9 plt.title('Median Price by ZIP Code')
         10 # Label every data point on the X-axis
         11 plt.xticks(range(len(median_prices)), median_prices['Zip Code'], ro
         12
         13
         14 plt.show()
         15

```



```
In [23]: 1 overall_median_price = kc_house_data_df['price'].median()
          2 print("Overall Median Price:", overall_median_price)
          3
```

Overall Median Price: 870000.0

I wanted to know what the Median Price of homes is in the Kings County dataset - It is \$870,000. Now I want to investigate which zip codes have a lot of homes for sale significantly below the median price for the county. I started by gathering the zipcodes that below the median price for the county.

```
In [24]: 1 low_median_prices = median_prices[median_prices['price'] < 870000]
```

I decided to create a new dataframe that consisted of all the zipcodes that were below the median price of \$870,000 for the county.

```
In [25]: 1 low_kc_house = kc_house_data_df[kc_house_data_df['Zip Code'].isin(l
```

```
In [26]: 1 low_kc_house.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 14431 entries, 7399300360 to 9557800100
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date                  14431 non-null  object
1   price                 14431 non-null  float64
2   bedrooms              14431 non-null  int64
3   bathrooms             14431 non-null  float64
4   sqft_living           14431 non-null  int64
5   sqft_lot              14431 non-null  int64
6   floors                14431 non-null  float64
7   waterfront            14431 non-null  object
8   greenbelt             14431 non-null  object
9   nuisance              14431 non-null  object
10  view                  14431 non-null  object
11  condition             14431 non-null  object
12  grade                 14431 non-null  object
13  heat_source           14431 non-null  object
14  ...                   ...
```



```
In [27]: 1 low_kc_house['grade'].value_counts()
```

```
Out[27]: 7 Average          6642
          8 Good           3991
          6 Low Average    2142
          9 Better         1051
          5 Fair           279
          10 Very Good      230
          11 Excellent       50
          4 Low             32
          3 Poor             8
          12 Luxury          5
          2 Substandard      1
          Name: grade, dtype: int64
```

Since the 'grade' column contained categorical data I wanted to change this to numerical data so it could be useful once I start modeling. I used the .map function and found information on this by typing map a function to a column in pandas and found this [solution](https://sparkbyexamples.com/pandas/pandas-map-function-explained/) (<https://sparkbyexamples.com/pandas/pandas-map-function-explained/>).

```
In [28]: 1 grade_mapping = {
          2     '2 Substandard': 2,
          3     '3 Poor': 3,
          4     '4 Low': 4,
          5     '5 Fair': 5,
          6     '6 Low Average': 6,
          7     '7 Average': 7,
          8     '8 Good': 8,
          9     '9 Better': 9,
         10     '10 Very Good': 10,
         11     '11 Excellent': 11,
         12     '12 Luxury': 12,
         13 }
         14
         15 low_kc_house['grade'] = low_kc_house['grade'].map(grade_mapping)
         16
         17
```

I wanted to check my new dataframes descriptive statistics and found that the median sqft_living of homes in these zipcodes was 1,810sqft. The median price of homes in this low median priced dataframe was \$674,950.

```
In [29]: 1 low_kc_house.describe()
```

```
Out[29]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	grade	sqft_abov
count	14431.00	14431.00	14431.00	14431.00	14431.00	14431.00	14431.00	14431.0
mean	739244.41	3.36	2.16	1924.63	16828.88	1.44	7.29	1672.7
std	375747.75	0.94	0.80	787.60	63370.81	0.52	0.97	726.4
min	27360.00	0.00	0.00	310.00	420.00	1.00	2.00	300.0
25%	542500.00	3.00	1.50	1340.00	5500.00	1.00	7.00	1146.0
50%	674950.00	3.00	2.00	1810.00	7680.00	1.00	7.00	1470.0
75%	838451.00	4.00	2.50	2380.00	10225.50	2.00	8.00	2080.0
max	8500000.00	10.00	9.50	8020.00	3253932.00	4.00	12.00	6490.0

I wanted to find the zipcodes that contained the most fair and low average grades because these are the homes that could easily return profits after minimal renovations were applied. The zipcodes of 98168, 98118, and 98146 had the highest amount.

```
In [30]: 1 # Filter the dataframe for "fair 5" or "low average 6" grades
2 fair_low_kc_house = low_kc_house[(low_kc_house['grade'] == 5) | (lo
3
4 # Group the data by Zip Code and count the occurrences
5 zipcode_counts = fair_low_kc_house['Zip Code'].value_counts()
6
7 # Display the result
8 print(zipcode_counts)
```

```
98168    177
98118    144
98146    143
98106    128
98002    122
98178    120
98126    109
98155    103
98022     99
98133     94
98056     91
98042     85
98001     85
98023     79
98166     77
98108     73
98058     70
98057     70
98125     65
98144     60
98038     59
98198     56
98070     55
98030     44
98188     41
98092     35
98010     29
98003     28
98055     26
98032     24
98031     20
98148     10
Name: Zip Code, dtype: int64
```

In [31]: 1 low_kc_house.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 14431 entries, 7399300360 to 9557800100
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   date                  14431 non-null  object
1   price                 14431 non-null  float64
2   bedrooms              14431 non-null  int64
3   bathrooms             14431 non-null  float64
4   sqft_living           14431 non-null  int64
5   sqft_lot              14431 non-null  int64
6   floors                14431 non-null  float64
7   waterfront            14431 non-null  object
8   greenbelt            14431 non-null  object
9   nuisance              14431 non-null  object
10  view                  14431 non-null  object
11  condition             14431 non-null  object
12  grade                 14431 non-null  int64
13  heat_source           14431 non-null  object
14  ...
```

In [32]: 1 *# Selecting specific columns*
2 selected_columns = ['price', 'bedrooms', 'bathrooms', 'sqft_living']
3 preview_df = low_kc_house[selected_columns]
4
5 *# Displaying the preview*
6 print(preview_df.head())

	price	bedrooms	bathrooms	sqft_living	grade	Zip Code
id						
7399300360	675000.00	4	1.00	1180	7	98055
8910500230	920000.00	5	2.50	2770	7	98133
1180000275	311000.00	6	2.00	2880	7	98178
1604601802	775000.00	3	3.00	2160	9	98118
2807100156	625000.00	2	1.00	1190	7	98133

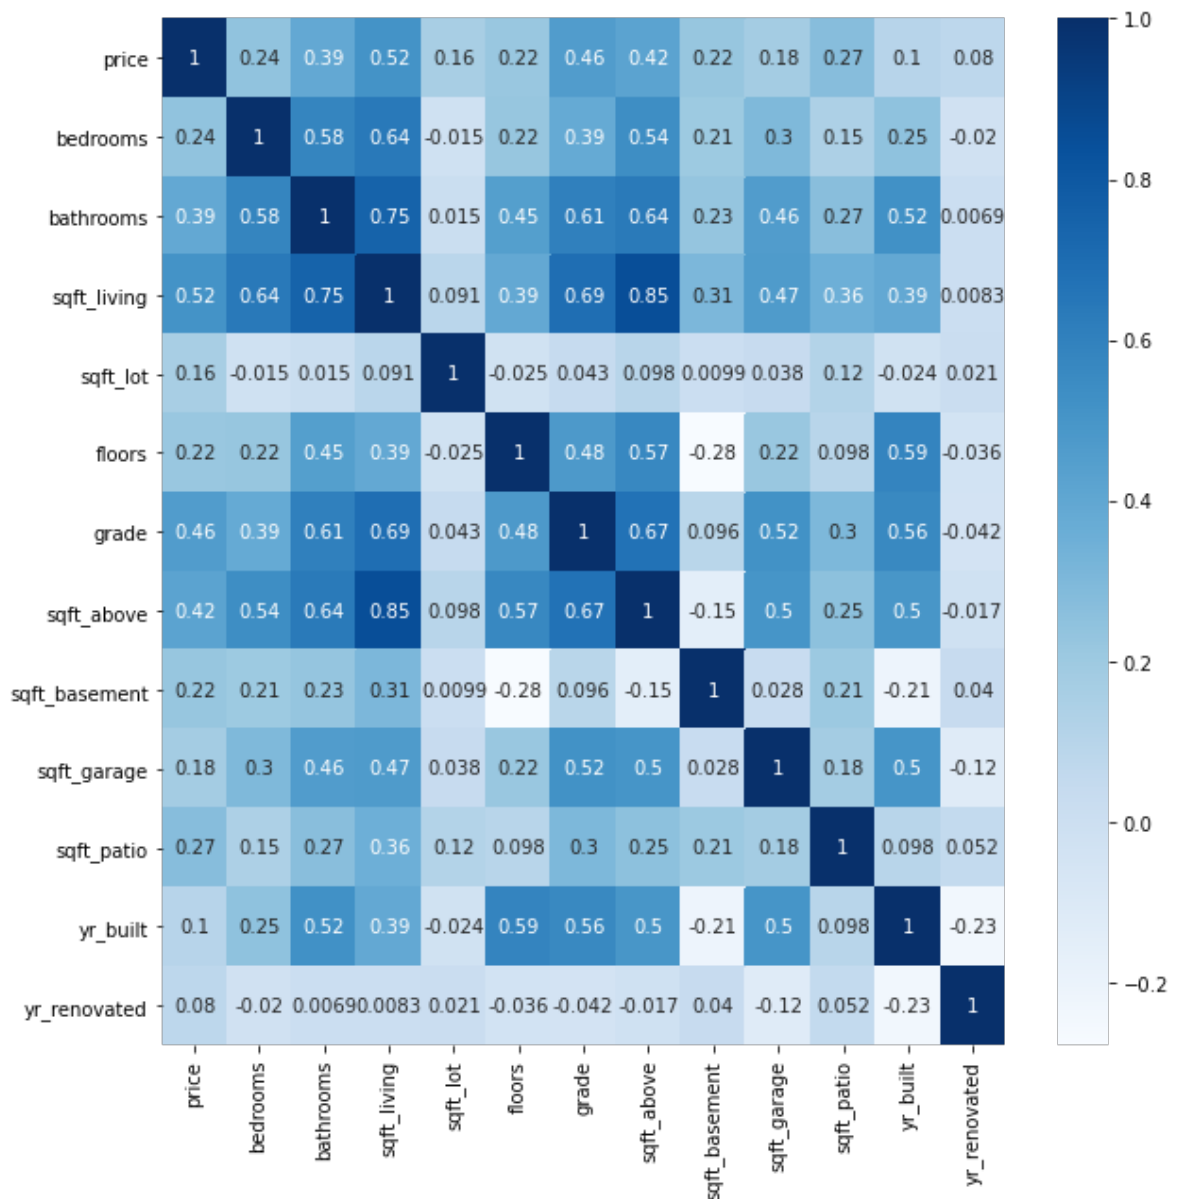
I dropped date because it was not a useful column of numeric values.

In [33]: 1 *# Drop columns 'date'*
2 low_kc_house.drop(columns=['date'], inplace=True)

Looking at another heatmap on the new dataframe and seeing sqft_living and grade are the highest correlated with price.

```
In [34]: 1 fig, ax = plt.subplots(figsize=(10,10))
          2 cor = low_kc_house.corr()
          3 sns.heatmap(cor, cmap="Blues", annot=True)
```

Out[34]: <AxesSubplot:>



Taking another look at what variables had the highest correlation with price.

```
In [35]: 1 low_kc_house.corr()['price']
```

```
Out[35]: price            1.00  
bedrooms          0.24  
bathrooms         0.39  
sqft_living       0.52  
sqft_lot          0.16  
floors            0.22  
grade             0.46  
sqft_above        0.42  
sqft_basement     0.22  
sqft_garage       0.18  
sqft_patio        0.27  
yr_built          0.10  
yr_renovated      0.08  
Name: price, dtype: float64
```

```
In [36]: 1 # Perform linear regression
2 coefficients = np.polyfit(low_kc_house['sqft_living'], low_kc_house
3 poly = np.poly1d(coefficients)
4
5 # Scatter plot
6 plt.figure(figsize=(10, 6))
7 plt.scatter(low_kc_house['sqft_living'], low_kc_house['price'], alp
8
9 # Plot the best-fit line
10 plt.plot(low_kc_house['sqft_living'], poly(low_kc_house['sqft_livin
11
12 # Title and labels
13 plt.title('Scatter Plot of sqft_living vs Price with Best-fit Line')
14 plt.xlabel('sqft_living')
15 plt.ylabel('Price (USD)')
16
17 # Set y-axis limit to show only positive values
18 plt.ylim(bottom=0)
19
20 # Convert y-axis ticks to represent prices in millions
21 plt.yticks(ticks=plt.yticks()[0], labels=['${:.0f}M'.format(x / 1e6
22
23 # Set axis intercept at (0, 0)
24 plt.axhline(0, color='black', linestyle='--', linewidth=1)
25 plt.axvline(0, color='black', linestyle='--', linewidth=1)
26
27 # Legend
28 plt.legend()
29
30 # Show the plot
31 plt.show()
```



Starting the train test split process. I entered my numerical features as my X variables and 'price' against my y variable.

```
In [37]: 1 # Define the features (X) and the target variable (y)
          2 X = low_kc_house[['sqft_living', 'bedrooms', 'bathrooms', 'floors', 'gr
          3 y = low_kc_house['price']
          4
```

```
In [38]: 1 # split the data into training and testing
          2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size
```

```
In [39]: 1 print(X_train.shape)
          2 print(X_test.shape)
          3 print(y_train.shape)
          4 print(y_test.shape)
```

```
(11544, 10)
(2887, 10)
(11544,)
(2887,)
```

```
In [40]: 1 X.dtypes
```

```
Out[40]: sqft_living      int64
          bedrooms      int64
          bathrooms    float64
          floors       float64
          grade         int64
          sqft_lot      int64
          sqft_above    int64
          sqft_garage   int64
          sqft_basement int64
          sqft_patio    int64
          dtype: object
```

I wanted to create a function so I could easily evaluate each of models with an r2 score, root mean squared error, and mean absolute error.


```
In [41]: 1 def evaluate_model(y_test, y_pred, lr):
2         # R-squared (R2)
3         r2 = r2_score(y_test, y_pred)
4
5         # Root Mean Squared Error (RMSE)
6         rmse = mean_squared_error(y_test, y_pred, squared=False)
7
8         # Mean Absolute Error (MAE)
9         mae = mean_absolute_error(y_test, y_pred)
10
11        # Intercept
12        intercept = lr.intercept_
13
14        # Printing the results
15        print("R2 score: ", r2)
16        print("Root Mean Squared Error: ", rmse)
17        print("Mean Absolute Error: ", mae)
18        #print("Intercept: ", intercept)
19
20        # Returning the results as a dictionary
21        results_model = {
22            'r2': r2,
23            'rmse': rmse,
24            'mae': mae,
25            #'intercept': intercept
26        }
27
28        return results_model
```

Modeling

Baseline Model

```
In [42]: 1 baseline_model = sm.OLS(endog=y_train, exog=sm.add_constant(X_train
2         baseline_model
```

```
Out[42]: <statsmodels.regression.linear_model.OLS at 0x7f6792cb7a30>
```

```
In [43]: 1 baseline_results = baseline_model.fit()
2         baseline_results
```

```
Out[43]: <statsmodels.regression.linear_model.RegressionResultsWrapper at 0x7f6792c2a880>
```

In [44]: 1 `print(baseline_results.summary())`

```

                                OLS Regression Results
=====
Dep. Variable:                  price    R-squared:
0.263
Model:                          OLS      Adj. R-squared:
0.263
Method:                        Least Squares    F-statistic:
4117.
Date:                          Mon, 27 Nov 2023    Prob (F-statistic):
0.00
Time:                          21:47:02    Log-Likelihood:
6271e+05
No. Observations:              11544    AIC:
3.254e+05
Df Residuals:                  11542    BIC:
3.254e+05
Df Model:                      1
Covariance Type:               nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
const          2.705e+05    7867.698     34.375     0.000     2.55e+05
2.86e+05
sqft_living    243.4649       3.794     64.164     0.000     236.027
250.903
=====
=====
Omnibus:                  13605.835    Durbin-Watson:
1.996
Prob(Omnibus):            0.000    Jarque-Bera (JB):      432
8467.790
Skew:                    5.844    Prob(JB):
0.00
Kurtosis:                97.140    Cond. No.
5.48e+03
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.48e+03. This might indicate that there are strong multicollinearity or other numerical problems.

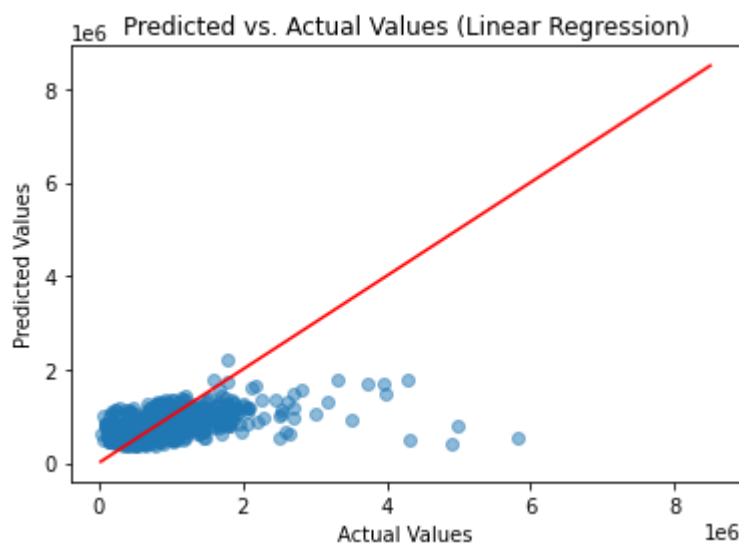
```
In [45]: 1 y_pred0 = baseline_results.predict(sm.add_constant(X_test['sqft_liv  
2 evaluate_model(y_test,y_pred0,baseline_model)
```

```
R2 score: 0.27453333546309444  
Root Mean Squared Error: 330222.30380844034  
Mean Absolute Error: 188870.09638627042
```

```
Out[45]: {'r2': 0.27453333546309444,  
          'rmse': 330222.30380844034,  
          'mae': 188870.09638627042}
```

The model explains about 26 percent of the variance which is fairly low. This is understandable because there is a lot more that determines the price of a house other than square footage. The model predicts that an increase of 1 sqft_living is associated with an increase of about 245 USD in price. The model is statistically significant with a p-value below the standard alpha of 0.05. The testing data is showing a r2 score (variance) of about 27 percent which means there is an ever so slight underfitting issue. The root mean squared error is very high at about \$330,000 which is substantial. The mean absolute error is also very high at about 188,000 USD. It should also be noted that the constant is at about 270,000 USD. This means that the model is being built on this hypothetical floor price which is also very high. We can conclude that although this model performs pretty well it should really not be used for any real world predictions.

```
In [46]: 1 # Plot the predicted vs. actual values  
2  
3 plt.scatter(y_test, y_pred0, alpha=0.5)  
4 plt.plot([y.min(), y.max()], [y.min(), y.max()], color="red")  
5 plt.xlabel("Actual Values")  
6 plt.ylabel("Predicted Values")  
7 plt.title("Predicted vs. Actual Values (Linear Regression)")  
8 plt.show()
```



Building a Multiple Linear Regression Model

```
In [47]: 1 X_train[['sqft_living', 'grade']]
```

Out[47]:

	sqft_living	grade
id		
7300400840	3120	9
3327900100	1800	7
9558020820	1900	8
3224049105	2420	7
4058801085	3110	8
...
7305300800	1200	6
520079082	3510	9
722069226	2990	7
87000114	860	5
1473120370	2790	9

11544 rows × 2 columns

Second Model

```
In [48]: 1 X = low_kc_house[['sqft_living', 'bedrooms', 'bathrooms', 'floors', 'gr  
2 X_second = low_kc_house[['sqft_living', 'grade']]  
3 y = low_kc_house['price']
```

```
In [49]: 1 MLR_model = sm.OLS(y_train, sm.add_constant(X_train[['sqft_living',
2 second_results = MLR_model.fit()
3 print(second_results.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:                  price    R-squared:
0.286
Model:                          OLS      Adj. R-squared:
0.285
Method:                        Least Squares    F-statistic:
2306.
Date:                          Mon, 27 Nov 2023    Prob (F-statistic):
0.00
Time:                          21:47:02    Log-Likelihood:
6253e+05                                -1.
No. Observations:              11544    AIC:
3.251e+05
Df Residuals:                  11541    BIC:
3.251e+05
Df Model:                      2
Covariance Type:               nonrobust
=====
=====
                                coef    std err          t      P>|t|      [0.025
0.975]
-----
const      -1.786e+05    2.48e+04    -7.217    0.000    -2.27e+05
-1.3e+05
sqft_living  175.8783        5.145     34.182    0.000     165.793
185.964
grade       7.945e+04    4159.037     19.103    0.000     7.13e+04
8.76e+04
=====
=====
Omnibus:                        13912.155    Durbin-Watson:
2.006
Prob(Omnibus):                  0.000    Jarque-Bera (JB):
0230.913                                489
Skew:                          6.061    Prob(JB):
0.00
Kurtosis:                      103.099    Cond. No.
1.77e+04
=====
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.77e+04. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [50]: 1 # Predict on the testing set
          2 # Y_preds are predicted values of the target variable based on the
          3 # Once we have those preds, we can compare them to the actual value
          4 # To evaluate the performance of the model
          5 y_pred1 = second_results.predict(sm.add_constant(X_test[['sqft_livi
```

```
In [51]: 1 evaluate_model(y_test,y_pred1,MLR_model)
```

```
R2 score:  0.29719144102045214
Root Mean Squared Error:  325024.5706205458
Mean Absolute Error:  185494.7150060409
```

```
Out[51]: {'r2': 0.29719144102045214,
          'rmse': 325024.5706205458,
          'mae': 185494.7150060409}
```

This model explains about 28.6% on the variance which is again fairly low but slightly better than the baseline model. The model predicts that 1 unit of increase in sqft_living would result in an increase of about 175 dollars in price. One unit of increase in grade predicts about a 79,000 dollar increase in price. The 2 features are statistically significant with low p-values. The testing r2 score is slightly higher at about 29.7% which means the model does have slight fit issues, most likely due to multicollinearity between the features. The root mean squared error is also very high at roughly 325,000 USD which is substantial but important to note that it is around 5,000 USD less than the baseline model. The mean absolute error is around 185,000 USD which is also very high but still roughly \$3,000 less than the baseline model. The constant is also lower at roughly 178,000 USD which is almost 100,000 USD less than the baseline model. Although this model does appear to be an improvement it should still not be used for real world predictions.

Third "Kitchen Sink" Model

```
In [52]: 1 # X = low_kc_house[['sqft_living', 'bedrooms', 'bathrooms', 'floors', '
          2 # y = low_kc_house['price']
```

In [53]:

```

1 MLR_model2 = sm.OLS(y_train, sm.add_constant(X_train))
2 third_results = MLR_model2.fit()
3 print(third_results.summary())

```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          price      R-squared:
0.332
Model:                  OLS      Adj. R-squared:
0.331
Method:                 Least Squares      F-statistic:
571.9
Date:                  Mon, 27 Nov 2023      Prob (F-statistic):
0.00
Time:                  21:47:02      Log-Likelihood:          -1.
6215e+05
No. Observations:      11544      AIC:
3.243e+05
Df Residuals:          11533      BIC:
3.244e+05
Df Model:              10
Covariance Type:       nonrobust
=====
=====
                                coef      std err          t      P>|t|      [0.025
0.975]
-----
const      -2.122e+05    2.74e+04    -7.747    0.000    -2.66e+05
-1.59e+05
sqft_living    128.7035     14.246     9.034    0.000     100.778
156.629
bedrooms     -4.185e+04    4128.568   -10.138    0.000    -4.99e+04
-3.38e+04
bathrooms     1.781e+04    5891.947     3.023    0.003     6263.663
2.94e+04
floors      -8314.9862    7389.614    -1.125    0.261    -2.28e+04
6169.912
grade         9.674e+04    4518.246    21.412    0.000     8.79e+04
1.06e+05
sqft_lot        0.6373       0.043    14.680    0.000         0.552
0.722
sqft_above     62.3448     14.338     4.348    0.000         34.241
90.449
sqft_garage   -183.5844     13.074   -14.042    0.000    -209.212
-157.956
sqft_basement   87.9452     10.093     8.714    0.000         68.162
107.728
sqft_patio    108.1793     14.186     7.626    0.000         80.373
135.986
=====
=====
Omnibus:          14195.026      Durbin-Watson:
2.003
Prob(Omnibus):    0.000      Jarque-Bera (JB):          601

```

```

2107.807
Skew:                    6.228    Prob(JB):
0.00
Kurtosis:                114.104    Cond. No.
6.71e+05
=====
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 6.71e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```

In [54]: 1 # Predict on the testing set
          2 # Y_preds are predicted values of the target variable based on the
          3 # Once we have those preds, we can compare them to the actual value
          4 # To evaluate the performance of the model
          5 y_pred2 = third_results.predict(sm.add_constant(X_test))

```

```

In [55]: 1 evaluate_model(y_test,y_pred2,MLR_model2)

```

```

R2 score:  0.35144510605774015
Root Mean Squared Error:  312227.42168256926
Mean Absolute Error:  176236.20491543153

```

```

Out[55]: {'r2': 0.35144510605774015,
          'rmse': 312227.42168256926,
          'mae': 176236.20491543153}

```

The "kitchen sink" model with all numeric features has a variance of about 33.2% which is again fairly low but also higher than the baseline model. It appears that all the features (with the exception of floors) are statistically significant with very low p-values. The testing r2 score is around 35% which means the model does have fit issues most likely due to multicollinearity between the features. The root mean squared error is very high at around 312,000, however it is has the lowest rmse of the 3 models. The feature Bedrooms has a negative correlation with price at roughly around 41,000 USD for a one unit increase. An example scenario where this would make sense is a new one or 2 bedroom condo in the heart of downtown would be more expensive than a 3 bedroom home in the middle of nowhere away from basic necessities like a grocery store. So we can see that there are some issues with the data. Practically you would think that a 4 bedroom home would be more valuable than a 2 bedroom home and in a lot of scenarios this would be true. This insight shows how truly important location must play a role in the price of a home.

Evaluation

The "kitchen sink" model is the best performing model. It has the highest variance of the 3 models but it is still fairly low at 33.2%. It has the lowest root mean squared error of the 3 models at roughly 312,000 which is still very large. For these reasons the model should not be

used for predictive measures. From the model we did learn that one unit of square footage is approximately equal to about 128 USD in price. A 1 unit increase in grade is approximately equal to 96,000 USD in price which seems high. It also shows that an additional bathroom is equal to about 17,000 USD in price. These are the features that seem to have the biggest impact on price.

Conclusion

The "kitchen sink" model is the best performing model. However this model should really not be used for real world predictions. We need to gather a lot more data (hundreds of thousands more entries) specifically in homes in the low median priced zipcode areas. Ideally looking for data that is not so multicollinear so that the regression models can be more useful/predictive. However in the meantime the low median priced zipcodes would be good areas to target. The top zipcode to focus on would be the '98168', '98118', and '98146' areas as it is showing the most homes with a 'fair' or 'low average' grade that in theory would return profits easily with minimal renovations.

Limitations

The main limitations of this dataset was that there was a lot of multicollinearity between the features of the homes. The features were more correlated with each other rather than the target variable (price). We could see this after looking at a heatmap (Pearson Correlation Coefficient Matrix) and after running each of the models with training/testing data. Although the Linear Regression models were helpful in getting an idea of how features correlate with price the models did not perform well enough to be predictive. Further the models were only built off around 14,000 entries of data gathering more (10x) along with using features that were not as colinear with each other would benefit the models predictive capabilities greatly.

Next Steps/Recommendations

We need a lot more data/records (hundreds of thousands) specifically in the zipcodes below the median price for homes in the county. Ideally trying to find features that do not have a lot of multicollinearity. Also obtaining data on schools in the district as well as crime reports in the county would be helpful as they would definitely have an impact on price. This kind of data would not be correlated with standard features of a home. Also gathering other economic data would be beneficial such as if there are major companies close by for work or other shopping plazas/restaurants nearby and how this would affect the price of homes.